

Maximum Likelihood Learning

Volodymyr Kuleshov

Cornell Tech


Lecture 4

Announcements

- Assignment 1 was released on Monday and is due on 2/13.
 - Please submit the assignment on Gradescope.
 - You can submit in teams of 1-2.
- The project proposal will be due on 2/22.
- I will release next week a sign-up sheet for student presentations.

The Task of Generative Modeling

Suppose we are given a training set of examples, e.g., images of dogs



Recall: A generative modeling algorithm has three components:

- **Model Family:** How do we parameterize a probability distribution over the data? What is the set of distributions we want to try to learn?
- **Learning Objective:** How do we quantify if a particular model in our family is a good fit?
- **Optimization:** How do we optimize the objective?

Step 1: How to represent $p(x)$ (last lecture). Step 2: how to learn it (today).

- ① Learning via Maximum Likelihood
 - The Learning Problem
 - The KL-Divergence as a Learning Objective
 - From KL-Divergence to Maximum Likelihood
- ② Maximizing the Likelihood from Data
 - Monte Carlo Estimation
 - Gradient Descent
- ③ Statistical Issues and the Bias/Variance Tradeoff

Running Example: A Generative Model for MNIST

Suppose we are given a dataset \mathcal{D} of handwritten digits (binarized MNIST)




- Each image has $n = 28 \times 28 = 784$ pixels. Each pixel can either be black (0) or white (1).
- **Our Goal:** Learn a probability distribution $p(x) = p(x_1, \dots, x_{784})$ over $x \in \{0, 1\}^{784}$ such that when $x \sim p(x)$, x looks like a digit
- Two step process:
 - 1 Parameterize a model family $\{p_\theta(x), \theta \in \Theta\}$ [This lecture]
 - 2 Search for model parameters θ based on training data \mathcal{D} [Next lecture]


Formalizing Assumptions on the Data

We start by making some standard assumptions on the data.

- We assume that the data comes from a probability distribution P_{data}
- Our dataset \mathcal{D} consists of m samples from P_{data}
 - Each sample is an assignment of values to (a subset of) the variables, e.g., $(X_{\text{bank}} = 1, X_{\text{dollar}} = 0, \dots, Y = 1)$ or pixel intensities.




$$\mathbf{x}_i \sim P_{\text{data}} \\ i = 1, 2, \dots, n$$




Formalizing the Task of Learning

Next, we want to define what “learning” means:

- We are also given a family of models \mathcal{M} , and our task is to learn some “good” model $P_\theta \in \mathcal{M}$ (i.e., in this family) that defines a distribution P_θ
 - For example, all Bayes nets with a given graph structure, for all possible choices of the CPD tables
 - For example, a FVSN for all possible choices of the logistic regression parameters. $\mathcal{M} = \{P_\theta, \theta \in \Theta\}$, θ = concatenation of all logistic regression coefficients
- The goal of learning is to return a model P_θ that precisely captures the distribution P_{data} from which our data was sampled



$$\mathbf{x}_i \sim P_{\text{data}} \\ i = 1, 2, \dots, n$$



What Defines a Good Model?

We may be interested in multiple tasks:


- ➊ **Density Estimation:** we are interested in the full distribution (for anomaly detection, missing value imputation, etc.)
- ➋ **Prediction:** we are using the distribution for supervised learning
 - Is this email spam or not?
 - Predict next frame in a video
- ➌ **Structure discovery:** we are interested in the model itself
 - How do some genes interact with each other?
 - What causes cancer?
- ➍ **Generation:** sample new data points that look “good”.

Our strategy will be to define an objective that balances among all these tasks.


Choosing a Learning Objective

We want to learn the full distribution so that later we can answer *any* probabilistic inference query

- In this setting we formulate a learning problem inspired by **density estimation**
- We want to construct P_θ as "close" as possible to P_{data} (recall we assume we are given a dataset \mathcal{D} of samples from P_{data})



$$\begin{aligned} \mathbf{x}_i &\sim P_{\text{data}} \\ i &= 1, 2, \dots, n \end{aligned}$$



- Our goal is to choose an objective $d(P_{\text{data}}, P_\theta)$ that captures "closeness" between $P_{\text{data}}, P_\theta$

A Learning Objective Based on the KL-Divergence

How should we measure "closeness" between distributions $P_{\text{data}}, P_{\theta}$?

- The **Kullback-Leibler divergence** (KL-divergence) between two distributions p and q is defined as

$$D(p\|q) = \sum_{\mathbf{x}} p(\mathbf{x}) \log \frac{p(\mathbf{x})}{q(\mathbf{x})}.$$


- $D(p\|q) \geq 0$ for all p, q , with equality if and only if $p = q$. Proof:

$$\mathbf{E}_{\mathbf{x} \sim p} \left[-\log \frac{q(\mathbf{x})}{p(\mathbf{x})} \right] \geq -\log \left(\mathbf{E}_{\mathbf{x} \sim p} \left[\frac{q(\mathbf{x})}{p(\mathbf{x})} \right] \right) = -\log \left(\sum_{\mathbf{x}} p(\mathbf{x}) \frac{q(\mathbf{x})}{p(\mathbf{x})} \right) = 0$$


- Hence, this is a sensible metric to quantify similarity of p and q
- Notice that KL-divergence is **asymmetric**, i.e., $D(p\|q) \neq D(q\|p)$

Learning as KL-Divergence Minimization

We are going to use the KL-Divergence to estimate the similarity between the distributions $P_{\text{data}}, P_{\theta}$



$\mathbf{x}_i \sim P_{\text{data}}$
 $i = 1, 2, \dots, n$



- The **KL-divergence** between $P_{\text{data}}, P_{\theta}$ is:

$$\mathbf{D}(P_{\text{data}} || P_{\theta}) = \mathbf{E}_{\mathbf{x} \sim P_{\text{data}}} \left[\log \left(\frac{P_{\text{data}}(\mathbf{x})}{P_{\theta}(\mathbf{x})} \right) \right] = \sum_{\mathbf{x}} P_{\text{data}}(\mathbf{x}) \log \frac{P_{\text{data}}(\mathbf{x})}{P_{\theta}(\mathbf{x})}$$

- $\mathbf{D}(P_{\text{data}} || P_{\theta}) = 0$ iff the two distributions are the same.

From KL-Divergence to Log-Likelihood

- We can simplify this somewhat:

$$\begin{aligned}\mathbf{D}(P_{\text{data}}||P_{\theta}) &= \mathbf{E}_{\mathbf{x} \sim P_{\text{data}}} \left[\log \left(\frac{P_{\text{data}}(\mathbf{x})}{P_{\theta}(\mathbf{x})} \right) \right] \\ &= \mathbf{E}_{\mathbf{x} \sim P_{\text{data}}} [\log P_{\text{data}}(\mathbf{x})] - \mathbf{E}_{\mathbf{x} \sim P_{\text{data}}} [\log P_{\theta}(\mathbf{x})]\end{aligned}$$

- The first term does not depend on P_{θ} .
- Then, *minimizing* KL divergence is equivalent to *maximizing* the **expected log-likelihood**

$$\arg \min_{P_{\theta}} \mathbf{D}(P_{\text{data}}||P_{\theta}) = \arg \min_{P_{\theta}} -\mathbf{E}_{\mathbf{x} \sim P_{\text{data}}} [\log P_{\theta}(\mathbf{x})] = \arg \max_{P_{\theta}} \mathbf{E}_{\mathbf{x} \sim P_{\text{data}}} [\log P_{\theta}(\mathbf{x})]$$

- Asks that P_{θ} assign high probability to instances sampled from P_{data} , so as to reflect the true distribution
- Because of log, samples \mathbf{x} where $P_{\theta}(\mathbf{x}) \approx 0$ weigh heavily in objective
- Although we can now compare models, since we are ignoring $\mathbf{H}(P_{\text{data}})$, we don't know how close we are to the optimum
- Problem: In general we do not know P_{data} .

- ① Learning via Maximum Likelihood
 - The Learning Problem
 - The KL-Divergence as a Learning Objective
 - From KL-Divergence to Maximum Likelihood
- ② **Maximizing the Likelihood from Data**
 - Monte Carlo Estimation
 - Gradient Descent
- ③ Statistical Issues and the Bias/Variance Tradeoff

Monte Carlo Estimation

To derive a learning objective, we will use *Monte Carlo estimation*:

- 1 Express the quantity of interest as the expected value of a random variable.

$$E_{x \sim P}[g(x)] = \sum_x g(x)P(x)$$

- 2 Generate T samples $\mathbf{x}^1, \dots, \mathbf{x}^T$ from the distribution P with respect to which the expectation was taken.
- 3 Estimate the expected value from the samples using:

$$\hat{g}(\mathbf{x}^1, \dots, \mathbf{x}^T) \triangleq \frac{1}{T} \sum_{t=1}^T g(\mathbf{x}^t)$$

where $\mathbf{x}^1, \dots, \mathbf{x}^T$ are independent samples from P . Note: \hat{g} is an *estimator* and a random variable.

Properties of the Monte Carlo Estimator

- **Unbiased:**

$$E_P[\hat{g}] = E_P[g(x)]$$

- **Convergence:** By law of large numbers

$$\hat{g} = \frac{1}{T} \sum_{t=1}^T g(x^t) \rightarrow E_P[g(x)] \text{ for } T \rightarrow \infty$$

- **Variance:**

$$V_P[\hat{g}] = V_P \left[\frac{1}{T} \sum_{t=1}^T g(x^t) \right] = \frac{V_P[g(x)]}{T}$$

Thus, variance of the estimator can be reduced by increasing the number of samples.

Empirical Maximum Likelihood

We can now use Monte Carlo to derive a practical learning objective:

- We approximate the expected log-likelihood

$$\mathbf{E}_{\mathbf{x} \sim P_{\text{data}}} [\log P_{\theta}(\mathbf{x})]$$

with the *empirical log-likelihood*:

$$\mathbf{E}_{\mathcal{D}} [\log P_{\theta}(\mathbf{x})] = \frac{1}{|\mathcal{D}|} \sum_{\mathbf{x} \in \mathcal{D}} \log P_{\theta}(\mathbf{x})$$

- **Maximum likelihood learning** is then:

$$\max_{P_{\theta}} \frac{1}{|\mathcal{D}|} \sum_{\mathbf{x} \in \mathcal{D}} \log P_{\theta}(\mathbf{x})$$

- Equivalently, we maximize probability of the data under model $P_{\theta}(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}) = \prod_{\mathbf{x} \in \mathcal{D}} P_{\theta}(\mathbf{x})$

An Example: Flipping a Biased Coin


Let's start with a simple example: flipping a biased coin

- Two outcomes: *heads* (H) and *tails* (T)
- Data set: Tosses of the biased coin, e.g., $\mathcal{D} = \{H, H, T, H, T\}$
- Assumption: the process is controlled by a probability distribution $P_{\text{data}}(x)$ where $x \in \{H, T\}$
- Class of models \mathcal{M} : Bernoulli distributions over $x \in \{H, T\}$.
- Example learning task: How should we choose $P_{\theta}(x)$ from \mathcal{M} if 60 out of 100 tosses are heads in \mathcal{D} ?

Maximum Likelihood Estimation for the Biased Coin

We represent our model: $P_\theta(x = H) = \theta$ and $P_\theta(x = T) = 1 - \theta$

- Example data: $\mathcal{D} = \{H, H, T, H, T\}$
- Likelihood of data = $\prod_i P_\theta(x_i) = \theta \cdot \theta \cdot (1 - \theta) \cdot \theta \cdot (1 - \theta)$



- Optimize for θ which makes \mathcal{D} most likely. What is the solution in this case?

MLE for the Coin Example: Analytical Derivation

Distribution: $P_{\theta}(x = H) = \theta$ and $P_{\theta}(x = T) = 1 - \theta$

- More generally, log-likelihood function

$$\begin{aligned}L(\theta) &= \theta^{\#heads} \cdot (1 - \theta)^{\#tails} \\ \log L(\theta) &= \log(\theta^{\#heads} \cdot (1 - \theta)^{\#tails}) \\ &= \#heads \cdot \log(\theta) + \#tails \cdot \log(1 - \theta)\end{aligned}$$

- MLE Goal: Find $\theta^* \in [0, 1]$ such that $\log L(\theta^*)$ is maximum.
- Differentiate the log-likelihood function with respect to θ and set the derivative to zero. We get:

$$\theta^* = \frac{\#heads}{\#heads + \#tails}$$

Extending the MLE Principle to a Bayesian Network

Given an autoregressive model with n variables and factorization

$$P_{\theta}(\mathbf{x}) = \prod_{i=1}^n p(x_i | pa(x_i); \theta_i)$$

Training data $\mathcal{D} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$. Maximum likelihood estimate of the parameters?

- Decomposition of Likelihood function

$$L(\theta, \mathcal{D}) = \sum_{j=1}^m \log P_{\theta}(\mathbf{x}^{(j)}) = \sum_{j=1}^m \sum_{i=1}^n \log p(x_i^{(j)} | pa(x_i)^{(j)}; \theta_i)$$

- Goal : maximize $\arg \max_{\theta} L(\theta, \mathcal{D}) = \arg \max_{\theta} \log L(\theta, \mathcal{D})$
- Each term is a normal conditional log-likelihood and can be optimized independently.
- For classical Bayes Net, conditionals are exponential families and have closed form solutions.

Extending the MLE Principle to a Neural Model

Given an autoregressive model with n variables and factorization

$$P_{\theta}(\mathbf{x}) = \prod_{i=1}^n p_{\text{neural}}(x_i | pa(x_i); \theta_i)$$

Training data $\mathcal{D} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$. Maximum likelihood estimate of the parameters?

- Decomposition of Likelihood function

$$L(\theta, \mathcal{D}) = \prod_{j=1}^m P_{\theta}(\mathbf{x}^{(j)}) = \prod_{j=1}^m \prod_{i=1}^n p_{\text{neural}}(x_i^{(j)} | pa(x_i)^{(j)}; \theta_i)$$

- Goal : maximize $\arg \max_{\theta} L(\theta, \mathcal{D}) = \arg \max_{\theta} \log L(\theta, \mathcal{D})$
- We no longer have a closed form solution!

MLE Learning: Gradient Descent

$$L(\theta, \mathcal{D}) = \prod_{j=1}^m P_{\theta}(\mathbf{x}^{(j)}) = \prod_{j=1}^m \prod_{i=1}^n p_{\text{neural}}(x_i^{(j)} | pa(x_i)^{(j)}; \theta_i)$$

Goal : maximize $\arg \max_{\theta} L(\theta, \mathcal{D}) = \arg \max_{\theta} \log L(\theta, \mathcal{D})$

$$\ell(\theta) = \log L(\theta, \mathcal{D}) = \sum_{j=1}^m \sum_{i=1}^n \log p_{\text{neural}}(x_i^{(j)} | pa(x_i)^{(j)}; \theta_i)$$

- 1 Initialize θ^0 at random
- 2 Compute $\nabla_{\theta} \ell(\theta)$ (by back propagation)
- 3 $\theta^{t+1} = \theta^t + \alpha_t \nabla_{\theta} \ell(\theta)$

Non-convex optimization problem, but often works well in practice

MLE Learning: Stochastic Gradient Descent

$$\ell(\theta) = \log L(\theta, \mathcal{D}) = \sum_{j=1}^m \sum_{i=1}^n \log p_{\text{neural}}(x_i^{(j)} | pa(x_i)^{(j)}; \theta_i)$$

- 1 Initialize θ^0 at random
- 2 Compute $\nabla_{\theta} \ell(\theta)$ (by back propagation)
- 3 $\theta^{t+1} = \theta^t + \alpha_t \nabla_{\theta} \ell(\theta)$

$$\nabla_{\theta} \ell(\theta) = \sum_{j=1}^m \sum_{i=1}^n \nabla_{\theta} \log p_{\text{neural}}(x_i^{(j)} | pa(x_i)^{(j)}; \theta_i)$$

What if $m = |\mathcal{D}|$ is huge?

$$\begin{aligned} \nabla_{\theta} \ell(\theta) &= m \sum_{j=1}^m \frac{1}{m} \sum_{i=1}^n \nabla_{\theta} \log p_{\text{neural}}(x_i^{(j)} | pa(x_i)^{(j)}; \theta_i) \\ &= m E_{x^{(j)} \sim \mathcal{D}} \left[\sum_{i=1}^n \nabla_{\theta} \log p_{\text{neural}}(x_i^{(j)} | pa(x_i)^{(j)}; \theta_i) \right] \end{aligned}$$

Monte Carlo: Sample $x^{(j)} \sim \mathcal{D}$; $\nabla_{\theta} \ell(\theta) \approx m \sum_{i=1}^n \nabla_{\theta} \log p_{\text{neural}}(x_i^{(j)} | pa(x_i)^{(j)}; \theta_i)$

Parallelization in Autoregressive Models

Our objective function is:

$$\ell(\theta) = \log L(\theta, \mathcal{D}) = \sum_{j=1}^m \sum_{i=1}^n \log p_{\text{neural}}(x_i^{(j)} | \text{pa}(x_i)^{(j)}; \theta_i)$$

If we use a recurrent neural network model, each term has the form

$$P(x_t | x_{1:t-1}) = P(x_t | h_{t-1}, x_{t-1}) \quad h_t = f(h_{t-1}, x_{t-1}).$$


Before we can evaluate and/or compute gradient, we need to process each term sequentially.

This is why we want feed-forward models like NADE, MADE, or WaveNet:

$$P(x_t | x_{1:t-1}) = P(x_t | x_{\text{neighborhood}}) = \text{conv}(x_{\text{neighborhood}})$$

Recall: WaveNet (Oord et al., 2016)

WaveNet is a state of the art model for speech:



WaveNet uses convolutions to parameterize a MADE-like masked autoencoder model mapping inputs x to model parameters \hat{x}_i .


- ① Learning via Maximum Likelihood
 - The Learning Problem
 - The KL-Divergence as a Learning Objective
 - From KL-Divergence to Maximum Likelihood
- ② Maximizing the Likelihood from Data
 - Monte Carlo Estimation
 - Gradient Descent
- ③ **Statistical Issues and the Bias/Variance Tradeoff**

Bias-Variance Trade-Off

- If the hypothesis space is very limited, it might not be able to represent P_{data} , even with unlimited data
 - This type of limitation is called **bias**, as the learning is limited on how close it can approximate the target distribution
- If we select a highly expressive hypothesis class, we might represent better the data
 - When we have small amount of data, multiple models can fit well, or even better than the true model. Moreover, small perturbations on \mathcal{D} will result in very different estimates
 - This limitation is call the **variance**.

Bias-Variance Trade-Off

- There is an inherent **bias-variance trade off** when selecting the hypothesis class. Error in learning due to both things: bias and variance.
- Hypothesis space: linear relationship
 - Does it fit well? Underfits
- Hypothesis space: high degree polynomial
 - Overfits
- Hypothesis space: low degree polynomial
 - Right tradeoff



How to Avoid Overfitting?

- Hard constraints, e.g. by selecting a less expressive hypothesis class:
 - Bayesian networks with at most d parents
 - Smaller neural networks with less parameters
 - Weight sharing
- Soft preference for “simpler” models: **Occam Razor**.
- Augment the objective function with **regularization**:

$$objective(\mathbf{x}, \mathcal{M}) = loss(\mathbf{x}, \mathcal{M}) + R(\mathcal{M})$$

- Evaluate generalization performance on a held-out validation set.
Log-likelihood should be similar on both training and validation set if there is no overfitting (as in discriminative modeling!)

Conditional Generative Models

- Suppose we want to generate a set of variables \mathbf{Y} given some others \mathbf{X} , e.g., text to speech
- We concentrate on modeling $p(\mathbf{Y}|\mathbf{X})$, and use a **conditional** loss function

$$-\log P_{\theta}(\mathbf{y} \mid \mathbf{x}).$$

- Since the loss function only depends on $P_{\theta}(\mathbf{y} \mid \mathbf{x})$, suffices to estimate the conditional distribution, not the joint



Input : image



Brown horse in
grass field

Output: caption

- For autoregressive models, it is easy to compute $p_\theta(x)$
 - Ideally, evaluate in parallel each conditional $\log p_{\text{neural}}(x_i^{(j)} | p_a(x_i)^{(j)}; \theta_i)$. Not like RNNs.
- Natural to train them via maximum likelihood
- Higher log-likelihood doesn't necessarily mean better looking samples
- Other ways of measuring similarity are possible (Generative Adversarial Networks, GANs)